



End-to-End Object Detection with Transformers

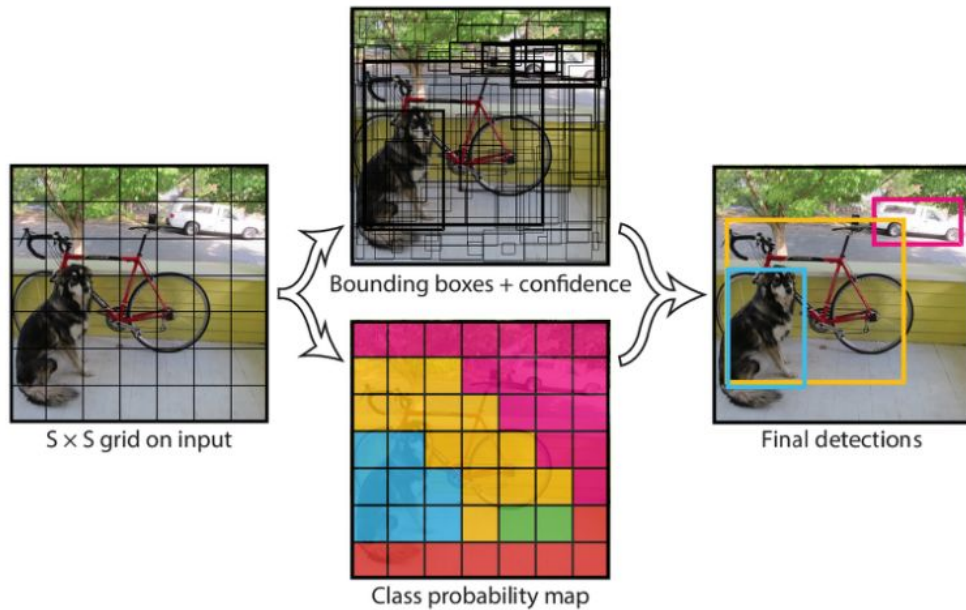

Nicolas Carion et al., ECCV 2020



Youngwoo Kim

DETR are a **set-based** global loss that forces unique predictions via **bipartite matching** and a **transformer encoder-decoder architecture**.


Conventional Object Detection



Anchor Box

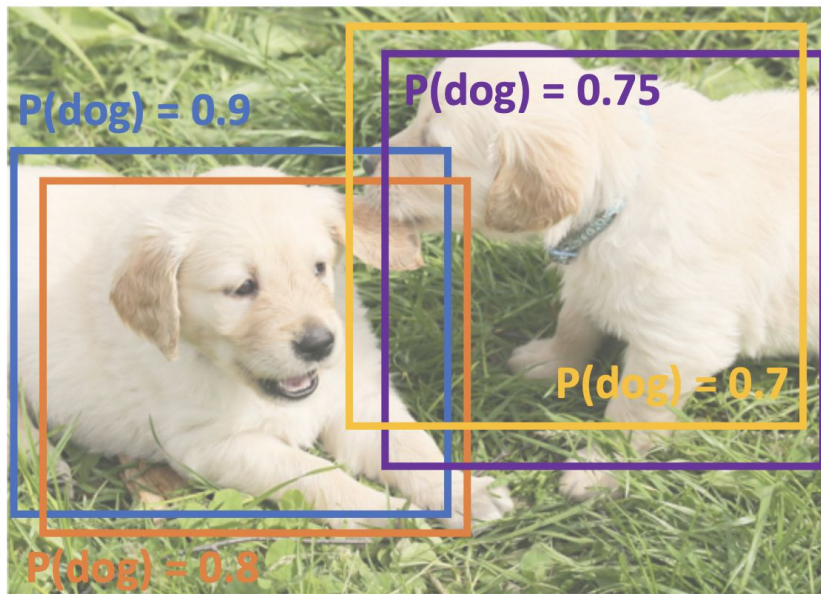


Input image
 $3 \times H \times W$

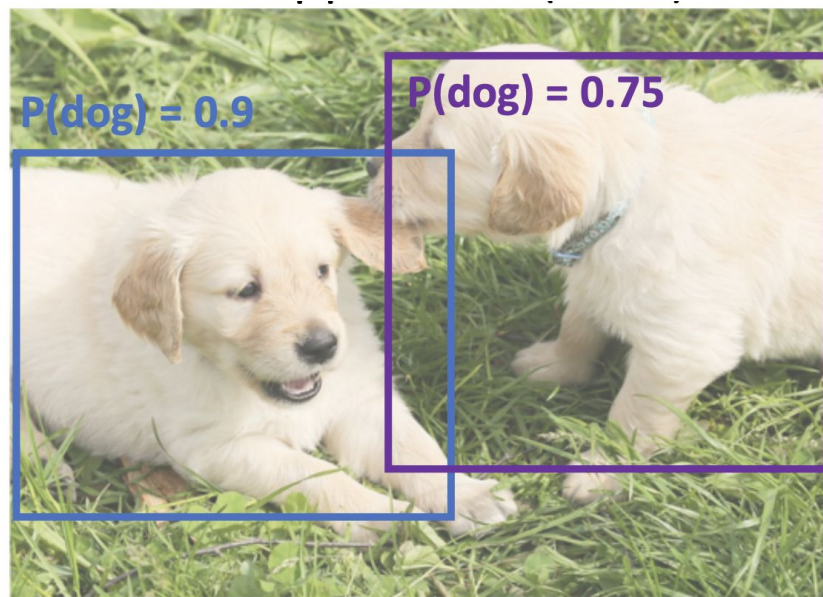


Divide image into grid
 7×7

Non Maximum Suppression


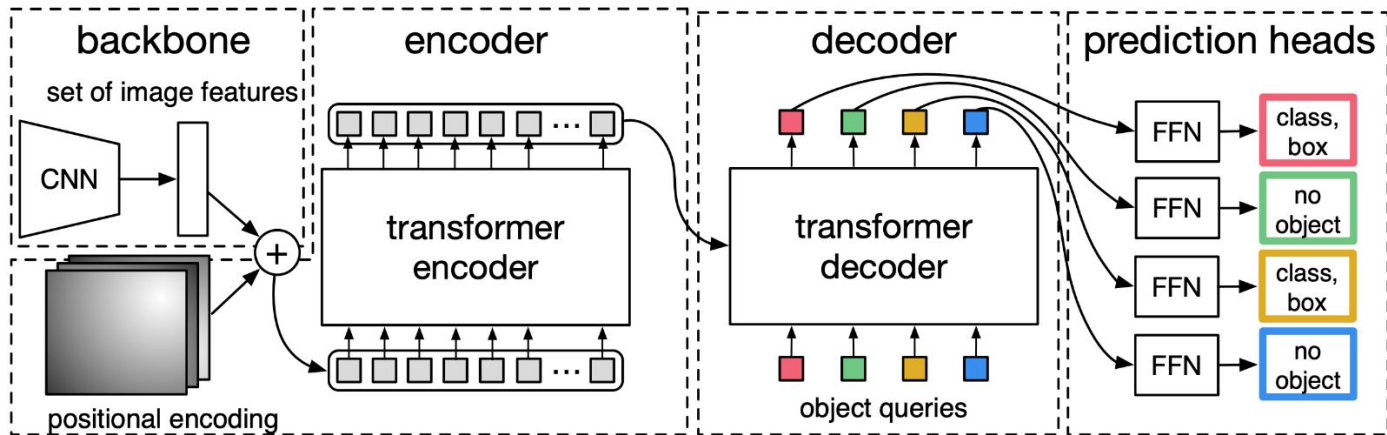


[Puppy image is CC0 Public Domain](#)



[Puppy image is CC0 Public Domain](#)

DETR



1. Bipartite matching
2. Transformer encoder-decoder


```



1  import torch
2  from torch import nn
3  from torchvision.models import resnet50
4
5  class DETR(nn.Module):
6
7      def __init__(self, num_classes, hidden_dim, nheads,
8                  num_encoder_layers, num_decoder_layers):
9          super().__init__()
10         # We take only convolutional layers from ResNet-50 model
11         self.backbone = nn.Sequential(*list(resnet50(pretrained=True).children())[:-2])
12         self.conv = nn.Conv2d(2048, hidden_dim, 1)
13         self.transformer = nn.Transformer(hidden_dim, nheads,
14                                         num_encoder_layers, num_decoder_layers)
15         self.linear_class = nn.Linear(hidden_dim, num_classes + 1)
16         self.linear_bbox = nn.Linear(hidden_dim, 4)
17         self.query_pos = nn.Parameter(torch.rand(100, hidden_dim))
18         self.row_embed = nn.Parameter(torch.rand(50, hidden_dim // 2))
19         self.col_embed = nn.Parameter(torch.rand(50, hidden_dim // 2))
20
21     def forward(self, inputs):
22         x = self.backbone(inputs)
23         h = self.conv(x)
24         H, W = h.shape[-2:]
25         pos = torch.cat([
26             self.col_embed[:W].unsqueeze(0).repeat(H, 1, 1),
27             self.row_embed[:H].unsqueeze(1).repeat(1, W, 1),
28         ], dim=-1).flatten(0, 1).unsqueeze(1)
29         h = self.transformer(pos + h.flatten(2).permute(2, 0, 1),
30                             self.query_pos.unsqueeze(1))
31         return self.linear_class(h), self.linear_bbox(h).sigmoid()
32
33 detr = DETR(num_classes=91, hidden_dim=256, nheads=8, num_encoder_layers=6, num_decoder_layers=6)
34 detr.eval()
35 inputs = torch.randn(1, 3, 800, 1200)
36 logits, bboxes = detr(inputs)

```

Find Optimal Matching

$$\hat{\sigma} = \arg \min_{\sigma \in \mathfrak{S}_N} \sum_i^N \mathcal{L}_{\text{match}}(y_i, \hat{y}_{\sigma(i)}),$$

Hungarian Algorithm





	1	2	∅	∅	∅
1	12	11	1	1	1
2	4	2	8	6	9
3	1	3	5	7	8
4	2	5	6	7	9
5	2	1	9	10	6

Permutation = [1, 2, 3, 4, 5]

Matching score = 12 + 2 + 5 + 7 + 6 = 32

Hungarian Algorithm



	1	2	\emptyset	\emptyset	\emptyset
1	12	11	1	1	1
2	4	2	8	5	9
3	1	3	5	7	8
4	2	5	6	7	4
5	2	1	9	10	6

Permutation = [3, 4, 1, 5, 2]

Matching score = 1 + 5 + 1 + 4 + 1 = 12

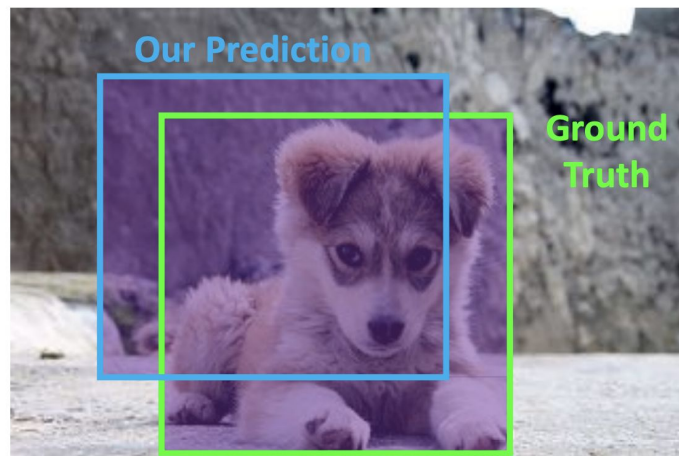

Bounding Box Loss

$$\mathcal{L}_{\text{box}}(b_i, \hat{b}_{\hat{\sigma}(i)}) = \lambda_{\text{iou}} \mathcal{L}_{\text{iou}}(b_i, \hat{b}_{\sigma(i)}) + \lambda_{\text{L1}} \|b_i - \hat{b}_{\sigma(i)}\|_1$$

Intersection over Union

$$\frac{\text{Area of Intersection}}{\text{Area of Union}}$$

$\text{IoU} > 0.5$ is “decent”,
 $\text{IoU} > 0.7$ is “pretty good”,
 $\text{IoU} > 0.9$ is “almost perfect”



Generalized IoU




Fig 5. GIoU loss

$$GIoU = IoU(b_{\sigma(i)}, \hat{b}) - \frac{|B(b_{\sigma(i)}, \hat{b})| \setminus b_{\sigma(i)} \cup \hat{b}_i}{|B(b_{\sigma(i)}, \hat{b})|}$$

$$\mathcal{L}_{iou}(b_{\sigma(i)}, \hat{b}) = 1 - GIoU$$

Compute Hungarian Loss

$$\mathcal{L}_{\text{Hungarian}}(y, \hat{y}) = \sum_{i=1}^N \left[-\log \hat{p}_{\hat{\sigma}(i)}(c_i) + \mathbb{1}_{\{c_i \neq \emptyset\}} \mathcal{L}_{\text{box}}(b_i, \hat{b}_{\hat{\sigma}(i)}) \right] ,$$

Backbone (Conventional CNN)

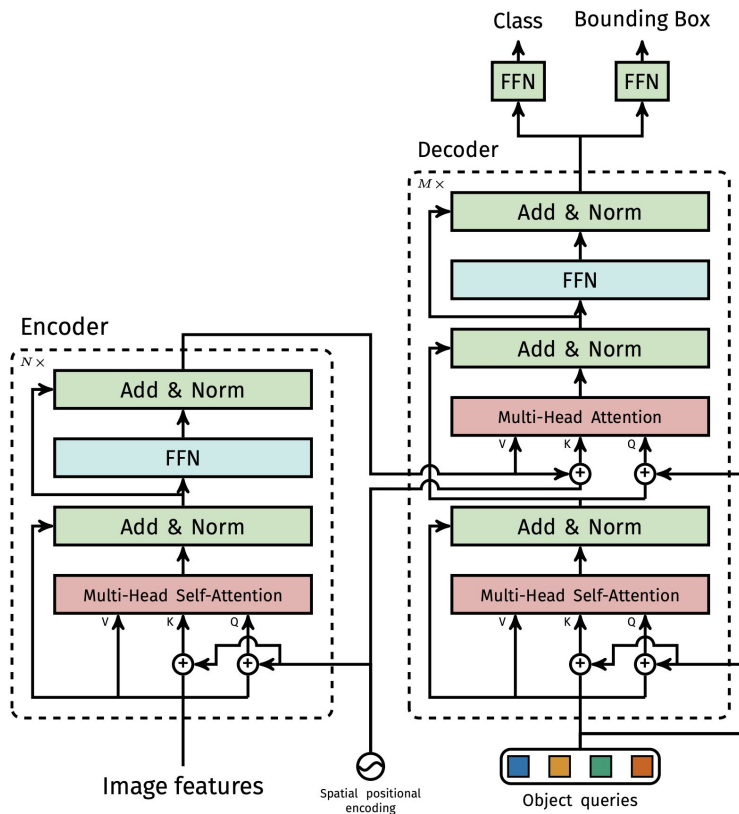
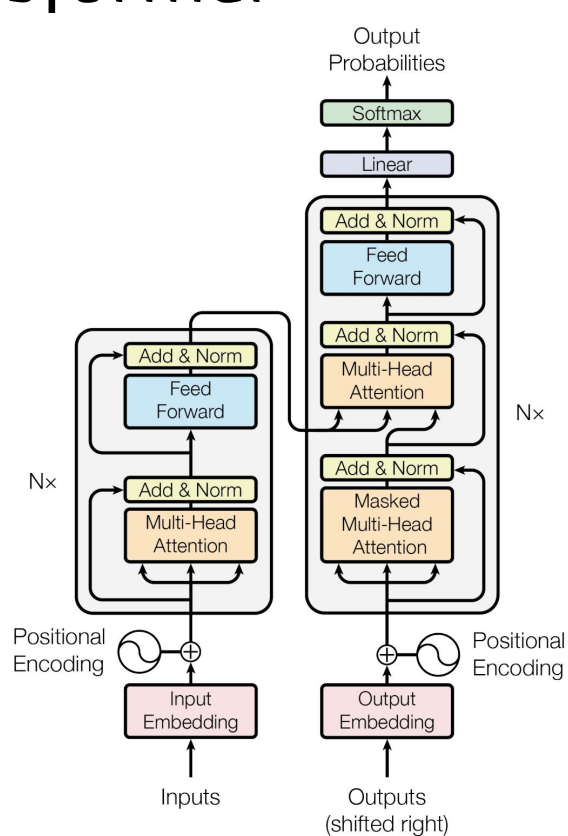
$$x_{\text{img}} \in \mathbb{R}^{3 \times H_0 \times W_0}$$



$$f \in \mathbb{R}^{C \times H \times W}$$

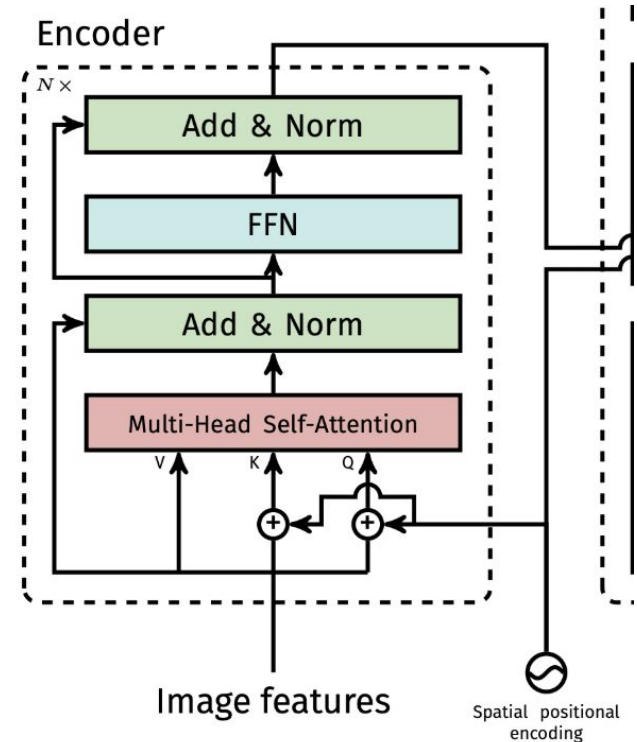
$$C = 2048 \quad H, W = \frac{H_0}{32}, \frac{W_0}{32}$$

Transformer




Transformer Encoder

- Reduces the channel dimension of the feature map from C to a smaller dimension $d = 256$
- Flatten feature map to $d \times HW$ feature map for transformer
- Fixed positional encodings are added to the input of each attention layer




Transformer Decoder

- decodes the N objects in parallel at each decoder layer, instead of using an autoregressive model, which predicts the output sequence one element at a time.
- Input embeddings, also known as object queries, are added to each attention layer



Prediction Feed Forward Networks

- Independently decodes object queries into bounding box coordinates and class labels
- The value of NNN is typically much larger than the actual number of objects



Faster R-CNN vs DETR


Model	GFLOPS/FPS	#params	AP	AP ₅₀	AP ₇₅	AP _S	AP _M	AP _L
Faster RCNN-DC5	320/16	166M	39.0	60.5	42.3	21.4	43.5	52.5
Faster RCNN-FPN	180/26	42M	40.2	61.0	43.8	24.2	43.5	52.0
Faster RCNN-R101-FPN	246/20	60M	42.0	62.5	45.9	25.2	45.6	54.6
Faster RCNN-DC5+	320/16	166M	41.1	61.4	44.3	22.9	45.9	55.0
Faster RCNN-FPN+	180/26	42M	42.0	62.1	45.5	26.6	45.4	53.4
Faster RCNN-R101-FPN+	246/20	60M	44.0	63.9	47.8	27.2	48.1	56.0
DETR	86/28	41M	42.0	62.4	44.2	20.5	45.8	61.1
DETR-DC5	187/12	41M	43.3	63.1	45.9	22.5	47.3	61.1
DETR-R101	152/20	60M	43.5	63.8	46.4	21.9	48.0	61.8
DETR-DC5-R101	253/10	60M	44.9	64.7	47.7	23.7	49.5	62.3

Number of encoder layers

- Without encoder layers, overall AP drops by 3.9
- AP on large object drops by 6.0


#layers	GFLOPS/FPS	#params	AP	AP ₅₀	AP _S	AP _M	AP _L
0	76/28	33.4M	36.7	57.4	16.8	39.6	54.2
3	81/25	37.4M	40.1	60.6	18.5	43.8	58.6
6	86/23	41.3M	40.6	61.6	19.9	44.3	60.2
12	95/20	49.2M	41.6	62.1	19.8	44.9	61.9

Visualizing encoder attention





Number of decoder layers

- As decoder layers increase, DETR performs good enough without NMS
- Both improved totalling into +8.2/9.5 AP between the first and the last layer



Visualizing decoder attention




Ablation - Loss

Table 4: Effect of loss components on AP. We train two models turning off ℓ_1 loss, and GIoU loss, and observe that ℓ_1 gives poor results on its own, but when combined with GIoU improves AP_M and AP_L . Our baseline (last row) combines both losses.

class	ℓ_1	GIoU	AP	Δ	AP_{50}	Δ	AP_S	AP_M	AP_L
✓	✓		35.8	-4.8	57.3	-4.4	13.7	39.8	57.9
✓		✓	39.9	-0.7	61.6	0	19.9	43.2	57.9
✓	✓	✓	40.6	-	61.6	-	19.9	44.3	60.2


Decoder output slot analysis

- DETR learns different specialization for each query slot

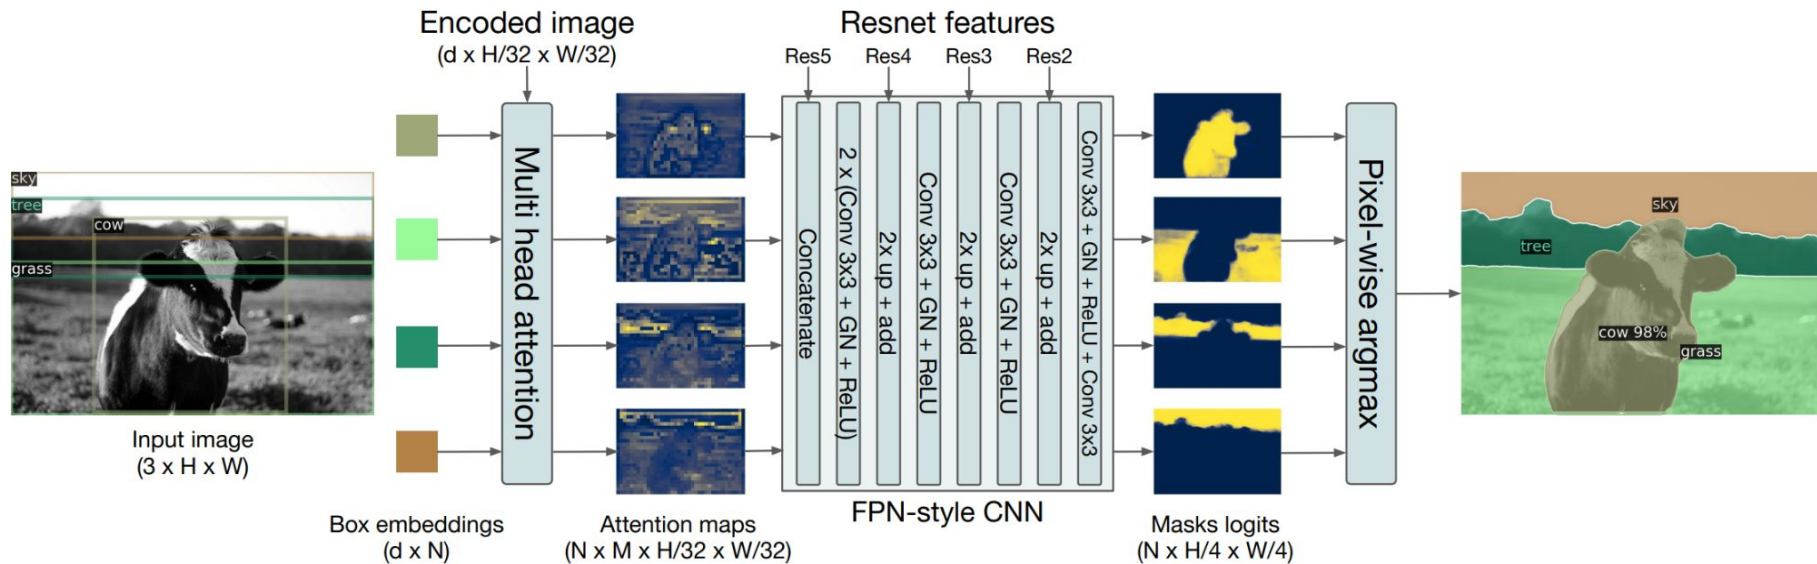


Generalization to unseen numbers of instances

- No image with more than 13 giraffes in the training set
- Able to detect all 24 giraffes
- Confirms that there is no strong class-specialization in each object query



DETR for panoptic segmentation







Table 5: Comparison with the state-of-the-art methods UPSNet [51] and Panoptic FPN [18] on the COCO val dataset We retrained PanopticFPN with the same data-augmentation as DETR, on a 18x schedule for fair comparison. UPSNet uses the 1x schedule, UPSNet-M is the version with multiscale test-time augmentations.

Model	Backbone	PQ	SQ	RQ	PQ th	SQ th	RQ th	PQ st	SQ st	RQ st	AP
PanopticFPN++	R50	42.4	79.3	51.6	49.2	82.4	58.8	32.3	74.8	40.6	37.7
UPSnet	R50	42.5	78.0	52.5	48.6	79.4	59.6	33.4	75.9	41.7	34.3
UPSnet-M	R50	43.0	79.1	52.8	48.9	79.7	59.7	34.1	78.2	42.3	34.3
PanopticFPN++	R101	44.1	79.5	53.3	51.0	83.2	60.6	33.6	74.0	42.1	39.7
DETR	R50	43.4	79.3	53.8	48.2	79.8	59.5	36.3	78.5	45.3	31.1
DETR-DC5	R50	44.6	79.8	55.0	49.4	80.5	60.6	37.3	78.7	46.5	31.9
DETR-R101	R101	45.1	79.9	55.5	50.5	80.9	61.7	37.0	78.5	46.0	33.0

Conclusion

- DETR is a new design for object detection systems based on transformers and bipartite matching loss for direct set prediction
- Despite some challenges, DETR successfully offers an end-to-end object detection system that does not require post-processing steps.
- Achieves comparable results to an optimized Faster R-CNN baseline on COCO dataset